

Govind_MoA_Machine_Learning_Final_Project

November 22, 2020

1 Mechanisms of Action (MoA) Prediction

Predicting multiple targets of the Mechanism of Action (MoA) response(s) of different samples (sig_id), given various inputs such as gene expression data and cell viability data.

1.1 Some of the important terms used in the headings of the tables are presented here:

g - : signifies gene expression data

c - : signifies cell expression data

cp_type : indicates samples treated with a compound (cp_vehicle) or with a control perturbation

NOTE: (samples with control perturbations don't have MoAs)

cp_time - treatment duration (24,48,72) Hours

cp_dose - Dosage - HIGH or LOW

```
[1]: # Importing the multi label stratified k-fold
      # cross validator

      from iterstrat.ml_stratifiers import MultilabelStratifiedKFold

      # Initial random imports
      import random
      import os
      import copy
      import warnings
      # warnings.filterwarnings('ignore')

      # Importing numpy
      import numpy as np

      # Importing pandas
      import pandas as pd

      # Importing matplotlib
      import matplotlib.pyplot as plt

      # Importing seaborn
```

```

import seaborn as sns

# Importing sklearn
from sklearn import preprocessing
from sklearn.metrics import log_loss
from sklearn.preprocessing import StandardScaler
from sklearn.decomposition import PCA
from sklearn.feature_selection import VarianceThreshold

# Importing pytorch
import torch
import torch.nn as nn
import torch.nn.functional as F
import torch.optim as optim

```

2 Using GPU if available

```

[2]: # using GPU if available
if torch.cuda.is_available():
    device_code = 'cuda'
else:
    device_code = 'cpu'

```

```

[3]: # setting the seed, so that every time the seed is started from the same number

def set_seed_characteristics(seed=55):
    # Setting a random seed value

    random.seed(seed)

    # for guaranteeing the reproducability of numbers by setting seed for NumPy
    np.random.seed(seed)

    # for setting the seed for cuda or cpu
    torch.manual_seed(seed)

    # To ensure that Pytorch doesnt just switch to the fastest possible
    ↪algorithm but
    # ensures that it selects a deterministic algorithm

    torch.backends.cudnn.deterministic = True

```

3 Reading the CSV Files

```
[4]: training_features = pd.read_csv('input/train_features.csv')
      # Reading the head rows and columns of train features
      training_features_head = training_features.head()

      training_targets_scored = pd.read_csv('input/train_targets_scored.csv')
      # Reading the head rows and columns of train targets scored
      training_targets_scored_head = training_targets_scored.head()

      testing_features = pd.read_csv('input/test_features.csv')
      # Reading the head rows and columns of train targets non-scored
      testing_features_head = testing_features.head()
```

```
[5]: # Printing the head - training features
      training_features_head
```

```
[5]:
```

	sig_id	cp_type	cp_time	cp_dose	g-0	g-1	g-2	g-3	\
0	id_000644bb2	trt_cp	24	D1	1.0620	0.5577	-0.2479	-0.6208	
1	id_000779bfc	trt_cp	72	D1	0.0743	0.4087	0.2991	0.0604	
2	id_000a6266a	trt_cp	48	D1	0.6280	0.5817	1.5540	-0.0764	
3	id_0015fd391	trt_cp	48	D1	-0.5138	-0.2491	-0.2656	0.5288	
4	id_001626bd3	trt_cp	72	D2	-0.3254	-0.4009	0.9700	0.6919	

	g-4	g-5	...	c-90	c-91	c-92	c-93	c-94	c-95	\
0	-0.1944	-1.0120	...	0.2862	0.2584	0.8076	0.5523	-0.1912	0.6584	
1	1.0190	0.5207	...	-0.4265	0.7543	0.4708	0.0230	0.2957	0.4899	
2	-0.0323	1.2390	...	-0.7250	-0.6297	0.6103	0.0223	-1.3240	-0.3174	
3	4.0620	-0.8095	...	-2.0990	-0.6441	-5.6300	-1.3780	-0.8632	-1.2880	
4	1.4180	-0.8244	...	0.0042	0.0048	0.6670	1.0690	0.5523	-0.3031	

	c-96	c-97	c-98	c-99
0	-0.3981	0.2139	0.3801	0.4176
1	0.1522	0.1241	0.6077	0.7371
2	-0.6417	-0.2187	-1.4080	0.6931
3	-1.6210	-0.8784	-0.3876	-0.8154
4	0.1094	0.2885	-0.3786	0.7125

[5 rows x 876 columns]

```
[6]: # Printing the head - train targets scored
      training_targets_scored_head
```

```
[6]:
```

	sig_id	5-alpha_reductase_inhibitor	11-beta-hsd1_inhibitor	\
0	id_000644bb2	0	0	
1	id_000779bfc	0	0	
2	id_000a6266a	0	0	

3	id_0015fd391	0	0
4	id_001626bd3	0	0

	acat_inhibitor	acetylcholine_receptor_agonist	\
0	0	0	
1	0	0	
2	0	0	
3	0	0	
4	0	0	

	acetylcholine_receptor_antagonist	acetylcholinesterase_inhibitor	\
0	0	0	
1	0	0	
2	0	0	
3	0	0	
4	0	0	

	adenosine_receptor_agonist	adenosine_receptor_antagonist	\
0	0	0	
1	0	0	
2	0	0	
3	0	0	
4	0	0	

	adenylyl_cyclase_activator	...	tropomyosin_receptor_kinase_inhibitor	\
0	0	...	0	
1	0	...	0	
2	0	...	0	
3	0	...	0	
4	0	...	0	

	trpv_agonist	trpv_antagonist	tubulin_inhibitor	\
0	0	0	0	
1	0	0	0	
2	0	0	0	
3	0	0	0	
4	0	0	0	

	tyrosine_kinase_inhibitor	ubiquitin_specific_protease_inhibitor	\
0	0	0	
1	0	0	
2	0	0	
3	0	0	
4	0	0	

	vegfr_inhibitor	vitamin_b	vitamin_d_receptor_agonist	wnt_inhibitor
0	0	0	0	0

1	0	0	0	0
2	0	0	0	0
3	0	0	0	0
4	0	0	0	0

[5 rows x 207 columns]

```
[7]: # Printing the head - test features
testing_features_head
```

```
[7]:      sig_id      cp_type  cp_time cp_dose      g-0      g-1      g-2      g-3  \
0  id_0004d9e33      trt_cp       24      D1 -0.5458  0.1306 -0.5135  0.4408
1  id_001897cda      trt_cp       72      D1 -0.1829  0.2320  1.2080 -0.4522
2  id_002429b5b  ctl_vehicle       24      D1  0.1852 -0.1404 -0.3911  0.1310
3  id_00276f245      trt_cp       24      D2  0.4828  0.1955  0.3825  0.4244
4  id_0027f1083      trt_cp       48      D1 -0.3979 -1.2680  1.9130  0.2057

      g-4      g-5  ...      c-90      c-91      c-92      c-93      c-94      c-95  \
0  1.5500 -0.1644  ...  0.0981  0.7978 -0.1430 -0.2067 -0.2303 -0.1193
1 -0.3652 -0.3319  ... -0.1190 -0.1852 -1.0310 -1.3670 -0.3690 -0.5382
2 -1.4380  0.2455  ... -0.2261  0.3370 -1.3840  0.8604 -1.9530 -1.0140
3 -0.5855 -1.2020  ...  0.1260  0.1570 -0.1784 -1.1200 -0.4325 -0.9005
4 -0.5864 -0.0166  ...  0.4965  0.7578 -0.1580  1.0510  0.5742  1.0900

      c-96      c-97      c-98      c-99
0  0.0210 -0.0502  0.1510 -0.7750
1  0.0359 -0.4764 -1.3810 -0.7300
2  0.8662  1.0160  0.4924 -0.1942
3  0.8131 -0.1305  0.5645 -0.5809
4 -0.2962 -0.5313  0.9931  1.8380
```

[5 rows x 876 columns]

4 Dataset classes, training and testing functions

```
[8]: # Pytorch data loader implementation of MoA dataset
class MoADataset:
    def __init__(self, features, targets):
        self.features = features
        self.targets = targets

    def __len__(self):
        return (self.features.shape[0])

    def __getitem__(self, idx):
```

```

train_tensor_dictionary = {
    'x' : torch.tensor(self.features[idx, :], dtype=torch.float),
    'y' : torch.tensor(self.targets[idx, :], dtype=torch.float)
}
return train_tensor_dictionary

# Pytorch data loader implementation of test dataset
class TestDataset:
    def __init__(self, features):
        self.features = features

    def __len__(self):
        return (self.features.shape[0])

    def __getitem__(self, idx):
        test_tensor_dictionary = {
            'x' : torch.tensor(self.features[idx, :], dtype=torch.float)
        }
        return test_tensor_dictionary

```

[10]: *# Pytorch model for the MoA determination*

```

class Model(nn.Module):

    # Instantiaing all the models before utilizing
    # them later in the forward function.
    def __init__(self, num_features, num_targets, hidden_size):

        # super keyword used to access data from the parent
        # pytorch.nn.Module class
        super(Model, self).__init__()
        # Applying batch normalization. This is done to standardize
        # the input for each mini batches and will help reduce the
        # number of epochs for which the training is done. This limits
        # the covariate shift (this is the value by which the hidden
        # layer values shift around) and allows to learn from a more
        # stable set of data. Sometimes, it also allows for a
        # higher learning rate. This is also used for regularization
        # and helps reduce over fitting. Generally, if batch
        # normalization is used, you can use a smaller dropout,
        # which in turn means that lesser layers can be lost
        # in every step.
        self.batch_normalization_1 = nn.BatchNorm1d(num_features)
        # For regularization purposes the dropout is set
        # This is done by setting a probablity. Random
        # neural networks are picked at a probablity, say p
        # or dropped at a probablity of 1-p. This is essential

```

```

# to prevent overfitting of the model and also reduces
# the computation time. A fully connected neural network, if
# run without dropout will start forming dependancies between
# each other and this can lead to over-fitting.
self.dropoutlayer_1 = nn.Dropout(0.2)
# nn.utils.weight_norm : This is weight normalization. Usually,
#                               faster than batch normalization
# nn.Linear : Applying linear transform to the incoming data
#               and creates a single layer feed forward network.
# input size : num_features
# output size : hidden_size
self.denselayer_1 = nn.utils.weight_norm(nn.Linear(num_features,
↪hidden_size))

self.batch_normalization_2 = nn.BatchNorm1d(hidden_size)
self.dropoutlayer_2 = nn.Dropout(0.2)
# input size : hidden_size
# output size : hidden_size
self.denselayer_2 = nn.utils.weight_norm(nn.Linear(hidden_size,
↪hidden_size))

self.batch_normalization_3 = nn.BatchNorm1d(hidden_size)
self.dropoutlayer_3 = nn.Dropout(0.1)
# input size : hidden_size
# output size : hidden_size
self.denselayer_3 = nn.utils.weight_norm(nn.Linear(hidden_size,
↪hidden_size))

self.batch_normalization_4 = nn.BatchNorm1d(hidden_size)
self.dropoutlayer_4 = nn.Dropout(0.1)
# input size : hidden_size
# output size : hidden_size
self.denselayer_4 = nn.utils.weight_norm(nn.Linear(hidden_size,
↪hidden_size))

self.batch_normalization_5 = nn.BatchNorm1d(hidden_size)
self.dropoutlayer_5 = nn.Dropout(0.1)
# input size : hidden_size
# output size : num_targets
self.denselayer_5 = nn.utils.weight_norm(nn.Linear(hidden_size,
↪num_targets))

# The forward function basically defines the model
def forward(self, forward_x):

    forward_x = self.batch_normalization_1(forward_x)

```

```

forward_x = self.dropoutlayer_1(forward_x)
forward_x = F.relu(self.denselayer_1(forward_x))

forward_x = self.batch_normalization_2(forward_x)
forward_x = self.dropoutlayer_2(forward_x)
forward_x = F.relu(self.denselayer_2(forward_x))

forward_x = self.batch_normalization_3(forward_x)
forward_x = self.dropoutlayer_3(forward_x)
forward_x = self.denselayer_3(forward_x)

forward_x = self.batch_normalization_4(forward_x)
forward_x = self.dropoutlayer_4(forward_x)
forward_x = self.denselayer_4(forward_x)

forward_x = self.batch_normalization_5(forward_x)
forward_x = self.dropoutlayer_5(forward_x)
forward_x = self.denselayer_5(forward_x)

return forward_x

```

```

[11]: # Function to train the model
def trainingFunction(model, optimizer, scheduler, lossFunction, trainloader, ↵
    ↪device_code):
    model.train()
    training_loss = 0
    for training_data in trainloader:
        optimizer.zero_grad()
        inputs, targets = training_data['x'].to(device_code), ↵
    ↪training_data['y'].to(device_code)
        outputs = model(inputs)
        loss = lossFunction(outputs, targets)
        loss.backward()
        optimizer.step()
        scheduler.step()
        training_loss += loss.item()
    training_loss /= len(trainloader)
    return training_loss

```

```

[12]: # Function to validate the model
def validationFunction(model, lossFunction, validationloader, device_code):
    model.eval()
    validation_loss = 0
    validation_predictions = []
    for validation_data in validationloader:
        inputs, targets = validation_data['x'].to(device_code), ↵
    ↪validation_data['y'].to(device_code)

```



```

        outputs = model(inputs)
        loss = lossFunction(outputs, targets)
        validation_loss += loss.item()
        validation_predictions.append(outputs.sigmoid().detach().cpu().numpy())
    validation_loss /= len(validationloader)
    validation_predictions = np.concatenate(validation_predictions)
    return validation_loss, validation_predictions

```

```

[13]: # Adding the inference function
def inferenceFunction(model, inferenceloader, device_code):
    model.eval()
    inferences = []
    for data in inferenceloader:
        inputs = data['x'].to(device_code)
        with torch.no_grad():
            outputs = model(inputs)
            inferences.append(outputs.sigmoid().detach().cpu().numpy())
    inferences = np.concatenate(inferences)
    return inferences

```

```

[14]: # Adding dummy inserts to the cp_time and cp_dose columns
# Usually done to categorical variables
def addDummies(data):
    dummy_data = pd.get_dummies(data, columns=['cp_time', 'cp_dose'])
    return dummy_data

```

5 Preparing the dataset

```

[15]: set_seed_characteristics(seed=55)

```

```

[16]: # Separating out the Gene expression Column and Cell Viability Column

gene_expression = [g for g in training_features.columns if g.startswith('g-')]
cell_viability = [c for c in training_features.columns if c.startswith('c-')]

```

```

[17]: # Since our dimensions are really high, we can resort to
# using PCA for dimensionality reduction, but is still able
# to capture the characteristics of the data.

# Now, this can be done by choosing a random dimension, and
# having the same random state as before. By doing this
# we observe that we do not encounter
# any 'nan' errors during training.

# Doing PCA for the Gene expression data

```

```

# can choose any random number here
random_pca_dimension_genes = 20

# Concatenating the training and test set
data = pd.concat([pd.DataFrame(training_features[gene_expression]), pd.
↳ DataFrame(testing_features[gene_expression])])

# Performing PCA and converting to a random_pca_dimension_genes number of
↳ columns
pca_genes = PCA(n_components = random_pca_dimension_genes, random_state=55)

# Fitting the PCA transform
data_pca = pca_genes.fit_transform(data[gene_expression])

# Splitting the training and test columns
train_pca_genes = data_pca[:training_features.shape[0]]
test_pca_genes = data_pca[-testing_features.shape[0]:]

# Converting training and testing into Pandas data frame shape
train_pca_genes = pd.DataFrame(train_pca_genes, columns=[f'pca_G-{i}' for i in
↳ range(random_pca_dimension_genes)])
test_pca_genes = pd.DataFrame(test_pca_genes, columns=[f'pca_G-{i}' for i in
↳ range(random_pca_dimension_genes)])

# Concatenating these back to the original features
training_features = pd.concat((training_features, train_pca_genes), axis=1)
testing_features = pd.concat((testing_features, test_pca_genes), axis=1)

```

[18]: # Doing PCA for the Cell Viability Data

```

# can choose any random number here
random_pca_dimension_cells = 32

# Concatenating the training and test set
data = pd.concat([pd.DataFrame(training_features[cell_viability]), pd.
↳ DataFrame(testing_features[cell_viability])])

# Performing PCA and converting to a random_pca_dimension_cells number of
↳ columns
pca_cells = PCA(n_components = random_pca_dimension_cells, random_state=55)

# Fitting the PCA transform
data_pca = pca_cells.fit_transform(data[cell_viability])

# Splitting the training and test columns

```

```

train_pca_cells = data_pca[:training_features.shape[0]]
test_pca_cells = data_pca[-testing_features.shape[0]:]

# Converting training and testing into Pandas data frame shape
train_pca_cells = pd.DataFrame(train_pca_cells, columns=[f'pca_C-{i}' for i in
    ↪range(random_pca_dimension_cells)])
test_pca_cells = pd.DataFrame(test_pca_cells, columns=[f'pca_C-{i}' for i in
    ↪range(random_pca_dimension_cells)])

# Concatenating these back to the original features
training_features = pd.concat((training_features, train_pca_cells), axis=1)
testing_features = pd.concat((testing_features, test_pca_cells), axis=1)

```

```

[19]: # Setting a desired threshold to calculate the VarianceThreshold.
# As per the math all the Features with a training-set variance
# lower than this threshold will be removed.
variancethreshold = VarianceThreshold(threshold=0.7)

# Combining training and test features to create a single dataset
combined_data = training_features.append(testing_features)

# Fits to the data, before transforming it
combined_data_transformed = variancethreshold.fit_transform(combined_data.iloc[
    ↪, 4:])

# Extracting the training and the testing data out of the
# transformed data
training_features_transformed = combined_data_transformed[ : training_features.
    ↪shape[0]]
testing_features_transformed = combined_data_transformed[-testing_features.
    ↪shape[0] : ]

```

```

[20]: # Extracting the training features in a suitable
# pandas dataset format and numbering the columns
# after the labels of 'sig_id', 'cp_type', 'cp_time', 'cp_dose'.
training_features = pd.
    ↪DataFrame(training_features[['sig_id', 'cp_type', 'cp_time', 'cp_dose']].values.
    ↪reshape(-1, 4), columns=['sig_id', 'cp_type', 'cp_time', 'cp_dose'])
training_features = pd.concat([training_features, pd.
    ↪DataFrame(training_features_transformed)], axis=1)

# Extracting the testing features in a suitable
# pandas dataset format and numbering the columns
# after the labels of 'sig_id', 'cp_type', 'cp_time', 'cp_dose'.

```

```

testing_features = pd.
↳DataFrame(testing_features[['sig_id','cp_type','cp_time','cp_dose']].values.
↳reshape(-1, 4), columns=['sig_id','cp_type','cp_time','cp_dose'])
testing_features = pd.concat([testing_features, pd.
↳DataFrame(testing_features_transformed)], axis=1)

```

```

[21]: # Merging the columns

train = training_features.merge(training_targets_scored, on='sig_id')

# Removing rows with cp_type as ctl_vehicle
# since control perturbations have no MoAs
# We are also manually setting the drop type as
# true because we do not want to include them back
# as a new column.

train = train[train['cp_type']!='ctl_vehicle'].reset_index(drop=True)

# Naturally, we have to get rid of them from the test dataset
# as well

test = testing_features[testing_features['cp_type']!='ctl_vehicle'].
↳reset_index(drop=True)

```

```

[22]: # Extracting the columns of the drugs that are sold from
# the train pandas dataframe

target = train[training_targets_scored.columns]

# Now that the ctl_vehicle drugs have been removed, we do not need
# cp_type. So we can go ahead and remove that columns as well.

train = train.drop('cp_type', axis=1)
test = test.drop('cp_type', axis=1)

# extracting the columns in the targets

target_columns = target.drop('sig_id', axis=1).columns.values.tolist()

```

6 Dataset preparation complete

```

[23]: # multilabel stratified K Fold import causes a small warning and we do not want
# to show that in the notebook.
warnings.filterwarnings('ignore')

```

```

folds = train.copy()
number_of_folds = 3

# creating a 3 fold multilabel stratified K Fold
multilabel_k_fold = MultilabelStratifiedKFold(n_splits = number_of_folds)

# Standard k fold splitting. Here we are splitting into number_of_folds folds

for fol, (train_folds, validation_folds) in enumerate(multilabel_k_fold.
    ↪split(X=train, y=target)):
    folds.loc[validation_folds, 'kfold'] = int(fol)

folds['kfold'] = folds['kfold'].astype(int)

# Isolating out the feature columns. This is done by first
# Isolating the columns that are not present in the target
# followed by extracting the columns except the sig_id and
# kfold.

feature_columns = [c for c in addDummies(folds).columns if c not in_
    ↪target_columns]
feature_columns = [c for c in feature_columns if c not in ['kfold','sig_id']]

```

7 Declaring the HyperParameters

```

[24]: BATCH_SIZE = 128
max_epochs = 16
# When training neural networks, it is common to use
# weight decay where after each update, the weights
# are multiplied by a factor slightly less than 1
weight_decay = 1e-5
# deciding the initial learning rate
# It controls how quickly or slowly a neural
# network model can learn a model or a problem.
lr = 1e-3
# Boolean to decide on stopping early when the
# validation_loss > best_loss
bool_early_stop = True
# steps to execute before early stopping
steps_early_stopping= 10
# number of features corresponding to the columns in the
# targets
num_features=len(feature_columns)
# number of targets corresponding to the columns in the
# features

```

```

num_targets=len(target_columns)
# in between neural network size
hidden_size=1024

```

8 Declaring the training functions and performing the training

```

[25]: # to plot validation loss
valid_loss_list = []
# to plot the training loss
train_loss_list = []
# to plot the best recorded loss
best_loss_list = []

```

```

[26]: def run_training(fold, seed):
    # declaring the list as global to plot validation loss
    global valid_loss_list
    # declaring the training loss list as global to plot
    # the training loss
    global train_loss_list
    # declaring the best loss list as global to plot the
    # best losses recorded
    global best_loss_list

    # setting the seed to start from the same number as
    # explained previously
    set_seed_characteristics(seed)

    # adding dummy variables to the training set
    train = addDummies(folds)

    # extracting the validating rows numbers for the
    # respective k fold values
    val_idx = train[train['kfold'] == fold].index

    # Dropping all the rows from the training set
    # that does not belong to this kth fold
    train_necessary_rows = train[train['kfold'] != fold].reset_index(drop=True)
    # Dropping all the rows from the validation set
    # that does not belong to this kth fold
    valid_necessary_rows = train[train['kfold'] == fold].reset_index(drop=True)

    # splitting the x and y values for training set
    train_features, train_targets = train_necessary_rows[feature_columns].
    ↪values, train_necessary_rows[target_columns].values
    # splitting the x and y values for test set

```

```

validation_features, validation_targets = ␣
↪valid_necessary_rows[feature_columns].values,␣
↪valid_necessary_rows[target_columns].values

# Converting the training data to standard pytorch
# dataset class format
train_dataset = MoADataset(train_features, train_targets)

# Converting the validation data to standard pytorch
# dataset class format
valid_dataset = MoADataset(validation_features, validation_targets)

# calling the pytorch data loading utility for the
# training set
trainloader = torch.utils.data.DataLoader(train_dataset,␣
↪batch_size=BATCH_SIZE, shuffle=True)
# calling the pytorch data loading utility for the
# validation set
validloader = torch.utils.data.DataLoader(valid_dataset,␣
↪batch_size=BATCH_SIZE, shuffle=False)

# Declaring the model and can be tuned here
# using the hyper parameters
model = Model(
    num_features=num_features,
    num_targets=num_targets,
    hidden_size=hidden_size,
)

# moving the model to GPU if available,
# else will run it on CPU itself
model.to(device_code)

# A standard optimizer. Adam optimizer is widely used
# because it combines the advantages of the Adaptive gradient
# algorithm and the root mean square propogation. Basically, it does
# not stick to one learning rate and adapts it to the problem.
# It is widely known to offer good results really fast.
optimizer = torch.optim.Adam(model.parameters(), lr=lr,␣
↪weight_decay=weight_decay)

# We use a learning rate scheduler to converge to the lowest
# loss faster. This is also seen to provide higher accuracy.
# This can be tuned.
# Some of the optimizers I tried here are
# optim.lr_scheduler.OneCycleLR
# optim.lr_scheduler.StepLR

```

```

# scheduler = optim.lr_scheduler.StepLR(optimizer, step_size=1, gamma=0.1)
scheduler = optim.lr_scheduler.OneCycleLR(optimizer=optimizer, pct_start=0.
→05, div_factor=1.5e3,
max_lr=1e-2, epochs=max_epochs,
→steps_per_epoch=len(trainloader))

# after research I saw that the Binary cross
# entropy loss with sigmoid layer works well
lossFunction = nn.BCEWithLogitsLoss()

# stops when the error starts increaseng. Setting the counter
# to track this
steps_before_early_stop = 0
# general out of fold array shape
out_of_fold = np.zeros((len(train), target.iloc[:, 1:].shape[1]))
# declaring a very high value as an
# initial loss for each kth fold
best_loss = np.inf

# looping through the epochs
for epoch in range(max_epochs):

    # training the model
    training_loss = trainingFunction(model, optimizer,scheduler,
→lossFunction, trainloader, device_code)
    print('epoch : ',epoch,'>> training_loss : ',training_loss)
    train_loss_list.append(training_loss)
    validation_loss, validation_predictions = validationFunction(model,
→lossFunction, validloader, device_code)
    print('epoch : ',epoch,'>> validation : ',validation_loss)
    valid_loss_list.append(validation_loss)

    # checking if the loss is decreasing
    if validation_loss < best_loss:
        best_loss = validation_loss
        best_loss_list.append(best_loss)
        # Updating the out of fold predictions
        out_of_fold[val_idx] = validation_predictions
        # saving the model and data for this kth fold
        torch.save(model.state_dict(), f"FOLD{fold}_.pth")

    # Handling the increasing loss by calling
    # early stopping
    elif(bool_early_stop == True):

        # breaks out of the loop when this happens

```



```

        steps_before_early_stop += 1
        if (steps_before_early_stop >= steps_early_stopping):
            break
# adding dummy variables to the test set
test_ = addDummies(test)

# extracting the x_test
x_test = test_[feature_columns].values
testdataset = TestDataset(x_test)
testloader = torch.utils.data.DataLoader(testdataset,
↳ batch_size=BATCH_SIZE, shuffle=False)

model = Model(num_features=num_features,num_targets=num_targets,
              hidden_size=hidden_size,
              )

# uploading the saved data for this kth fold
model.load_state_dict(torch.load(f"FOLD{fold}_pth"))
# again uploading the model to GPU, if available
model.to(device_code)

predictions = np.zeros((len(test_), target.iloc[:, 1:].shape[1]))
# evaluates the model
predictions = inferenceFunction(model, testloader, device_code)

return out_of_fold, predictions

```

```

[27]: def executeKFold(number_of_folds, seed):
    # standard size for the out of fold predictions
    out_of_fold = np.zeros((len(train), len(target_columns)))
    # same size for all of the predictions
    predictions = np.zeros((len(test), len(target_columns)))

    for each_k_fold in range(number_of_folds):
        print('Fold Number : ', each_k_fold)
        out_of_fold_, pred_ = run_training(each_k_fold, seed)

        # adding all the predictions
        predictions += pred_ / number_of_folds
        # adding all the out of fold predictions
        out_of_fold += out_of_fold_
        print("-----")

    k_th_prediction = predictions
    return out_of_fold, k_th_prediction

```

```
[28]: # setting a standard seed number
SEED = [55]
# general out of fold array shape
out_of_fold = np.zeros((len(train), len(target_columns)))
# general predictions array shape
predictions = np.zeros((len(test), len(target_columns)))

# for seed in SEED:
out_of_fold_, predictions_ = executeKFold(number_of_folds, SEED[0])
out_of_fold += out_of_fold_ / len(SEED)
predictions += predictions_ / len(SEED)

train[target_columns] = out_of_fold
test[target_columns] = predictions
```

```
Fold Number : 0
epoch : 0 >> training_loss : 0.2788736069979875
epoch : 0 >> validation : 0.021623456240471066
epoch : 1 >> training_loss : 0.020091968256494272
epoch : 1 >> validation : 0.01922281719101914
epoch : 2 >> training_loss : 0.019972266159627747
epoch : 2 >> validation : 0.01886744482506966
epoch : 3 >> training_loss : 0.019043210343174313
epoch : 3 >> validation : 0.018996690303601068
epoch : 4 >> training_loss : 0.018774466394730235
epoch : 4 >> validation : 0.018413769980442935
epoch : 5 >> training_loss : 0.01869301253362842
epoch : 5 >> validation : 0.01853232747264977
epoch : 6 >> training_loss : 0.01865103912094365
epoch : 6 >> validation : 0.018168202376571196
epoch : 7 >> training_loss : 0.018572327146387617
epoch : 7 >> validation : 0.01888334468520921
epoch : 8 >> training_loss : 0.018362585228422414
epoch : 8 >> validation : 0.018195184943234098
epoch : 9 >> training_loss : 0.018180621188619864
epoch : 9 >> validation : 0.01839715341941036
epoch : 10 >> training_loss : 0.017992399317090926
epoch : 10 >> validation : 0.017737267185644858
epoch : 11 >> training_loss : 0.01772677853865468
epoch : 11 >> validation : 0.017560527595723498
epoch : 12 >> training_loss : 0.017529506098641002
epoch : 12 >> validation : 0.01734824900932867
epoch : 13 >> training_loss : 0.017258823679193208
epoch : 13 >> validation : 0.01717964628839801
epoch : 14 >> training_loss : 0.01701616107121758
epoch : 14 >> validation : 0.01715383331837325
epoch : 15 >> training_loss : 0.016877019105722074
```

epoch : 15 >> validation : 0.01710507392497926

Fold Number : 1

epoch : 0 >> training_loss : 0.2791394925635794
epoch : 0 >> validation : 0.021501707748092454
epoch : 1 >> training_loss : 0.02014602518276028
epoch : 1 >> validation : 0.019920575047104525
epoch : 2 >> training_loss : 0.019292888320658518
epoch : 2 >> validation : 0.020500779746036077
epoch : 3 >> training_loss : 0.01898241023654523
epoch : 3 >> validation : 0.019408218819519568
epoch : 4 >> training_loss : 0.01884553834957921
epoch : 4 >> validation : 0.018803634969838733
epoch : 5 >> training_loss : 0.018704059833417767
epoch : 5 >> validation : 0.01929768891041649
epoch : 6 >> training_loss : 0.018495980099491452
epoch : 6 >> validation : 0.018345469753418504
epoch : 7 >> training_loss : 0.018502545138092144
epoch : 7 >> validation : 0.0183179613748758
epoch : 8 >> training_loss : 0.01830926013543554
epoch : 8 >> validation : 0.018034464246111697
epoch : 9 >> training_loss : 0.01814349767468546
epoch : 9 >> validation : 0.018158412088864838
epoch : 10 >> training_loss : 0.01796842050617156
epoch : 10 >> validation : 0.017744970556091647
epoch : 11 >> training_loss : 0.017735546859710113
epoch : 11 >> validation : 0.01759479866074077
epoch : 12 >> training_loss : 0.01754570741854284
epoch : 12 >> validation : 0.017433413593419666
epoch : 13 >> training_loss : 0.017299547661905702
epoch : 13 >> validation : 0.017249190232491697
epoch : 14 >> training_loss : 0.017085145907881467
epoch : 14 >> validation : 0.017144549509574628
epoch : 15 >> training_loss : 0.01694276431656402
epoch : 15 >> validation : 0.01715218401032275

Fold Number : 2

epoch : 0 >> training_loss : 0.28062024895587695
epoch : 0 >> validation : 0.021831682176682455
epoch : 1 >> training_loss : 0.019925031827195832
epoch : 1 >> validation : 0.02110225957786215
epoch : 2 >> training_loss : 0.019255759826172954
epoch : 2 >> validation : 0.01923344059494035
epoch : 3 >> training_loss : 0.01897560269774302
epoch : 3 >> validation : 0.01963105608291667
epoch : 4 >> training_loss : 0.01888560847905667
epoch : 4 >> validation : 0.018987211848384346
epoch : 5 >> training_loss : 0.018881849539668663

```

epoch : 5 >> validation : 0.019087364800788206
epoch : 6 >> training_loss : 0.01880092287193174
epoch : 6 >> validation : 0.018451588421032346
epoch : 7 >> training_loss : 0.01856426177141459
epoch : 7 >> validation : 0.018233612982620453
epoch : 8 >> training_loss : 0.018412961470692053
epoch : 8 >> validation : 0.018018708034450638
epoch : 9 >> training_loss : 0.018320308147889118
epoch : 9 >> validation : 0.01789587644603232
epoch : 10 >> training_loss : 0.018012238172409326
epoch : 10 >> validation : 0.017706910565752406
epoch : 11 >> training_loss : 0.01781968524598557
epoch : 11 >> validation : 0.01754140906870879
epoch : 12 >> training_loss : 0.017585085947876392
epoch : 12 >> validation : 0.017338143388644374
epoch : 13 >> training_loss : 0.01736602078637351
epoch : 13 >> validation : 0.017160253085452933
epoch : 14 >> training_loss : 0.01719682110554498
epoch : 14 >> validation : 0.017085314712262357
epoch : 15 >> training_loss : 0.017056497793806635
epoch : 15 >> validation : 0.0171179612135065
-----

```

9 Evaluating the logarithmic loss function applied to each drug-MoA annotation pair.

```

[29]: validation_df = training_targets_scored.drop(columns=target_columns).
      ↪merge(train[['sig_id']+target_columns], on='sig_id', how='left').fillna(0)
      # True target values
      true_target = training_targets_scored[target_columns].values
      # Predicted target values
      predicted_target = validation_df[target_columns].values
      cross_validation_score = 0

      # Now we can calculate the cross entropy loss

      for i in range(len(target_columns)):
          cross_validation_score_target = log_loss(true_target[:, i],
      ↪predicted_target[:, i])
          cross_validation_score += cross_validation_score_target / target.shape[1]


      print(" The Cross Validation loss is :>> ", cross_validation_score)

```

The Cross Validation loss is :>> 0.015734850054263057


10 Plotting the Validation loss for each fold

```
[39]: plt.plot(valid_loss_list)
plt.title('Validation Loss')
plt.savefig('valid_loss_list.png')
```



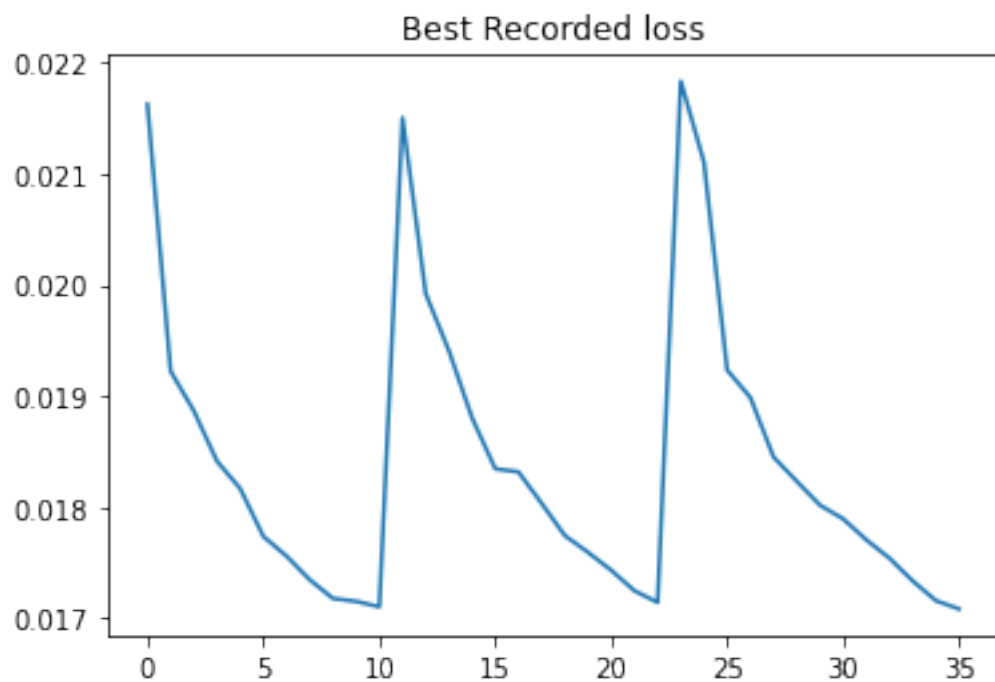
11 Plotting the Training loss for each fold

```
[40]: plt.plot(train_loss_list)
plt.title('Training loss')
plt.savefig('train_loss_list.png')
```



12 Plotting the Best recorded loss for each fold

```
[41]: plt.plot(best_loss_list)
      plt.title('Best Recorded loss')
      plt.savefig('best_loss_list.png')
```



[]: